Unified Reinforcement and Imitation Learning for Vision-Language Models

Byung-Kwan Lee† NVIDIA, KAIST byungkwanl@nvidia.com Ryo Hachiuma NVIDIA rhachiuma@nvidia.com Yong Man Ro KAIST ymro@kaist.ac.kr

Yu-Chiang Frank Wang NVIDIA, National Taiwan University frankwang@nvidia.com Yueh-Hua Wu NVIDIA krisw@nvidia.com

Abstract

Vision-Language Models (VLMs) have achieved remarkable progress, yet their large scale often renders them impractical for resource-constrained environments. This paper introduces Unified Reinforcement and Imitation Learning (RIL), a novel and efficient training algorithm designed to create powerful, lightweight VLMs. RIL distinctively combines the strengths of reinforcement learning with adversarial imitation learning. This enables smaller student VLMs not only to mimic the sophisticated text generation of large teacher models but also to systematically improve their generative capabilities through reinforcement signals. Key to our imitation framework is an LLM-based discriminator that adeptly distinguishes between student and teacher outputs, complemented by guidance from multiple large teacher VLMs to ensure diverse learning. This unified learning strategy, leveraging both reinforcement and imitation, empowers student models to achieve significant performance gains, making them competitive with leading closed-source VLMs. Extensive experiments on diverse vision-language benchmarks demonstrate that RIL significantly narrows the performance gap with state-of-the-art open- and closed-source VLMs and, in several instances, surpasses them. [Project Page]




Figure 1: Showing the performance improvements (%) of Qwen2.5-VL-7B [1] across vision-language evaluation benchmarks for AI2D [2], MathVista [3], MM-Vet [4], MMMU [5], BLINK [6], and the average scores for 14 evaluation benchmarks used in Table 1. Note that, we conduct RL on GRPO [7] and advanced GRPO, Dr.GRPO [8], with only answer rewards from LLM-as-a-Judge [9] (see Algorithm 1), and we present RIL based on similarity rewards from single or multi large teacher VLMs and simultaneously answer rewards (see Algorithm 2).

[†] Work Done during Internship.




Figure 2: Comparing **RIL**-applied VLMs based on multi large VLMs with diverse open- and closed-source VLMs, under average performance of numerous vision-language evaluation benchmarks: AI2D [2], ChartQA [19], MathVista [3], MMB [20], MM-Vet [4], MMMU [5], MMMU-Pro [21], MMStar [22], BLINK [6], SEED [23], SEED2+ [24], and RealWorldQA (RWQA).

1 Introduction

The pursuit of artificial general intelligence (AGI) has gained momentum, with much of the initial progress driven by instruction-tuned large language models (LLMs) [10–15]. Vision-Language Models (VLMs) [16, 1, 17, 18] represent a significant step forward, extending the LLM framework to integrate visual data and thereby achieve multimodal understanding. Their ability to generate insightful, visually-grounded textual responses that increasingly emulate human-like comprehension has made VLMs a focal point of considerable attention. Recognizing their transformative potential, numerous organizations are heavily investing to push the boundaries of what VLMs can achieve.

Historically, the enhancements in vision-language performance have primarily relied on traditional strategies such as scaling up model sizes [15, 25, 26] and expanding visual instruction tuning datasets [27–30, 18]. More recent advancements have explored architectural modifications—for instance, directly altering model structures [31, 32] or integrating auxiliary vision [33, 27, 34] and reasoning modules [35]. Alongside these efforts, the 'think-answer' paradigm [36–42] has gained prominence following DeepSeek-R1's [7] application of reinforcement learning (RL) via GRPO [43]. However, such architectural changes and the verbose 'think' responses preceding answers can significantly increase inference latency and computational memory requirements. These challenges render many powerful VLMs impractical for resource-limited settings like smartphones and augmented reality (AR) devices. As a result, there is a growing imperative within the research community to develop VLM designs that achieve a compelling balance between strong vision-language capabilities, low inference latency, and a lightweight model footprint.

To overcome these challenges, we introduce an approach for developing high-performing, efficient VLMs that avoids architectural modifications and the need for lengthy 'think' responses. Our method, termed Unified Reinforcement and Imitation Learning (RIL), is an efficient training algorithm integrating principles from both GRPO [43] and the GAIL framework [44]. The central objective of RIL is to enable smaller VLMs (e.g., 7B models) to effectively mimic the text generation style of significantly larger VLMs (e.g., 72B models). Prior works [45, 46] have argued that natural language response-based distillation is more effective than high-dimensional feature distillation, emphasizing the importance of utilizing the language head—referred to as the verbalization effect. Motivated by this insight, we began by leveraging the similarity in natural language responses between teacher and student models. Drawing inspiration from how models articulate responses differently, RIL employs an LLM-based discriminator trained to distinguish between the text responses generated by the student (small VLMs) and the teacher (large VLMs). This discriminator's output provides a similarity reward signal, guiding the student VLMs, which act as a generator, to produce responses increasingly akin to those of the teacher. To maintain training stability and prevent the generator or discriminator from overpowering the other, a common issue known as the balance problem [47, 48], the discriminator's architecture initially mirrors that of the generator.

However, the practical implementation of RIL for VLMs entails several specific challenges. Firstly, relying on continuous discriminator scores (ranging from zero to one) to assess similarity to large VLM outputs can introduce ambiguity into the learning signal. To ensure a clearer and more decisive reward, and drawing inspiration from prior works that binarize answer rewards [7, 43], we convert the discriminator's similarity score into a binary value. Secondly, the discriminator, by design, focuses on stylistic similarity and does not inherently verify factual correctness against ground truth

SYSTEM: You are a discriminator assistant.
You are provided with Response in <response> </response> tag.
Your task is to determine whether the Response was generated by teacher or student model.
Return 0 if the Response was generated by the teacher model.
Return 1 if the Response was generated by the student model.

USER: <response>{response}</response>
Determine the source of the Response.
Provide your answer as one of the following: 0 (Teacher model) 1 (Student model)

Figure 3: Input prompt for Discriminator

answers, which could otherwise lead to performance degradation. We mitigate this by incorporating an LLM-as-a-Judge [9] to provide a separate binary reward signal based on the accuracy of the generated response. A third challenge arises during the GRPO [43] phase: if the student VLMs, due to their more limited knowledge, fail to generate any correct responses for a given task, they struggle to identify effective learning pathways. To provide clearer guidance, especially when the student VLM lacks sufficient domain-specific knowledge compared to larger models, we incorporate text responses from both student and teacher VLMs during this GRPO training step. This strategy not only stabilizes RIL training but also creates opportunities for the student VLM to potentially surpass the performance of larger teachers.

Several key aspects of RIL contribute to its effectiveness. Notably, we observe that sourcing text responses from **multiple large teacher VLMs** significantly boosts student performance beyond what a single teacher can achieve, owing to the richer diversity of responses and the consequently more robust discriminator. This enriched learning environment is further enhanced during GRPO [43] step: by incorporating text responses from both student and teacher VLMs, we provide the student with clearer exemplars for reaching correct answers. RIL also demonstrates particular synergy with distillation-based VLMs; these models show marked performance improvements because their inherent feature alignment through distillation complements RIL's objectives. Beyond these learning dynamics, a crucial practical advantage of RIL is its inference efficiency: trained models do not require an explicit 'think' phase before generating answers, thereby maintaining the same inference speed as the original student VLM.

Furthermore, RIL's reliance on LLM-as-a-Judge [9] for accuracy evaluation offers substantial advantages in applicability. Unlike methods such as DeepSeek-R1 [7] and its variants [39, 49, 41, 37], which often use conventional answer parsing for reward computation and are thus limited to domains with predefined metrics (e.g., math or image grounding), RIL transcends these limitations. For example, conventional parsing struggles with open-ended general visual questions, such as "How to cook this dish?" based on a food image, or "Summarize this movie" from its poster. In addition, the parsing poses a challenge to flexibly compare different types of answers, if a response is "The answer is twenty percent" but its ground truth is "20". The versatile evaluation power of LLM-as-a-Judge [9] enables RIL's successful application to a broad spectrum of visual question answering tasks, fostering strong performance in areas like recognition, OCR, common-sense reasoning, mathematical problem-solving, and chart/document understanding. Finally, because RIL operates purely on generated text responses, it is **agnostic to the specific image embedding strategies or language tokenizers** used by the student or teacher VLMs, ensuring broad compatibility.

To validate our approach, we conduct extensive experiments across diverse vision-language benchmarks. The results compellingly demonstrate that RIL not only significantly narrows the performance gap to state-of-the-art open- and closed-source VLMs but, in several instances, surpasses them. The primary contributions of this work are:

- Unified Reinforcement and Imitation Learning (RIL): We introduce RIL, an efficient and novel training framework for VLMs. It empowers smaller models to emulate the text generation patterns of significantly larger VLMs, leading to substantial enhancements in their overall performance.
- **Broad Applicability and Flexibility**: RIL exhibits wide applicability, functioning effectively with diverse VLMs irrespective of their underlying image embedding strategies or language tokenizers. Furthermore, it preserves the original inference speed by avoiding lengthy intermediate reasoning steps and capably addresses general visual question answering tasks through its integration with LLM-as-a-Judge.

```
SYSTEM: You are an evaluation assistant.

Question is in <question> </question> tag.

Ground Truth is in <ground truth> </ground truth> tag.

Generated Text in <generated text> </generated text> tag.

After reading the Question, compare the Generated Text against the Ground Truth summary:

- If the Generated Text fully and correctly captures the core point → 1

- If it is incorrect or irrelevant → 0

Output the evaluation score (0 or 1) after giving a brief explanation

- The evaluation score should be wrapped in <answer> </answer> tags

USER: <question>{question}</a>

'question> <a href="mailto:qround truth">qround truth</a>

'ground truth>{ground truth}

'generated text> Provide the evaluation score after giving a brief explanation.

The evaluation score should be wrapped in <answer> </answer> tag.
```

Figure 4: Input prompt for LLM-as-a-Judge [9]

• Extensive Validation and Significant Performance Gains: Our comprehensive experimental results consistently show that RIL delivers substantial performance improvements across a variety of vision-language benchmarks. These gains position RIL-trained models as highly competitive with, and sometimes superior to, existing state-of-the-art VLMs.

2 Related Works

Evolution of Vision-Language Models (VLMs). Conventional remedies for advancing VLMs have largely centered on scaling up model sizes and expanding datasets to push performance boundaries. Consequently, promising VLMs have emerged, for example, closed-source VLMs: GPT-40 [50], Gemini [51], and Claude-3.5 Sonnet [52], and open-source VLMs: Molmo-72B [53], LLaVA-OneVision-72B [28], NVLM-72B [54], Qwen2.5-VL-72B [1], and InternVL3-78B [18] have followed this paradigm, incorporating large-scale language models such as Qwen2/Qwen2.5-72B [13]. However, the sheer size and computational demands of these models present significant barriers to deployment in resource-constrained environments such as mobile and embedded devices. To address these challenges, CoLLaVO [33] and MoAI [34] utilize computer vision models directly for visual capability, and Mini-Gemini [27], MoVA [55], and Eagle [56] employ multiple vision encoders such as CLIP [57], ConvNeXt [58], DINO-v2 [59], and SAM [60]. In parallel, Meteor [35] explores the efficient way of learning complex reasoning abilities, and TroL [31] and Phantom [32] investigate propagation modification for how we can embed vision-language knowledge as much as possible despite using the same architectures. More recently, since DeepSeek-R1 [7] introduces think-answer process and shows its dramatic performance improvements by using reinforcement learning (RL) such as GRPO [43], many variant models for VLMs has presented LMM-R1 [61], Vision-R1 [37], R1-V [42], OpenVLThinker [62], R1-OneVision [63], R1-Zero [64], and MM-Eureka [65]. They employ think-answer process to VLMs in specific areas such as math problems and object counting tasks. In addition, NoisyRollout [40] injects noise into clean images for distortion and train VLMs to strongly improve visual robustness, thereby understanding visual properties more than before.

Imitation Learning (IL). It originates from the concept of learning expert behavior in robotics. Generative adversarial imitation learning (GAIL) [44] is a foundational framework in this domain. It employs a generator and a discriminator where the generator (e.g., small student VLMs) attempts to replicate expert behavior (e.g., large teacher VLMs) by producing outputs such as actions or trajectories (e.g., text responses) that are indistinguishable from those generated by experts. Meanwhile, the discriminator aims to distinguish between the generator and the expert. By framing imitation learning (IL) as a minimax game, GAIL [44] leverages adversarial training to align the behavior of the generator with that of the expert without explicitly defining a reward function. Instead, GAIL [44] only relies on discriminator scores to evaluate how the generator similarly produces outputs of an expert. This approach has demonstrated strong performance in complex, high-dimensional tasks by directly learning from expert data, making it a widely adopted strategy in IL [66]. Based on its effectiveness, we integrate reinforcement and imitation learning (RIL) for VLMs with four key modifications: (1) combining GRPO [43] and GAIL [44] framework with explicitly reward design, (2) making the output scores of the discriminator binary to stabilize training, (3) incorporating LLM-as-a-Judge [9] to evaluate answer rewards that assess whether the generated text responses from

Table 1: Comparing the performances by using answer rewards from LLM-as-a-Judge [9] with purely RL on GRPO [7] and advanced GRPO [8], and by using similarity and answer rewards from RIL on GAIL [44], under numerous evaluation benchmarks: AI2D [2], ChartQA [19], MathVista [3], MMB/MMB^{CN} [20], MM-Vet [4], MM-Vet-v2 [67], MMMU [5], MMMU-Pro [21], MMStar [22], BLINK [6], SEED [23], SEED2+ [24], and RealWorldQA (RWQA). Note that, for RIL, we set large VLMs as the largest version of same small VLMs, such as Qwen2.5-VL-7B \leftarrow Qwen2.5-VL-72B and InternVL3-8B \leftarrow InternVL3-78B.

VLMs	AI2D	ChartQA	MathVista	MMB	MMB ^{CN}	MM-Vet	MM-Vet-v2	MMMU	MMMU-Pro	MMStar	BLINK	SEED	SEED2+	RWQA
Qwen2.5-VL-7B	83.9	87.3	67.8	83.5	83.4	71.8	63.7	55.0	38.3	63.9	56.4	77.0	70.4	68.5
w. RL (GRPO)	84.1	88.8	69.0	83.8	83.9	72.8	63.9	56.2	40.3	65.2	58.9	77.4	70.6	69.4
w. RL (Dr.GRPO)	84.5	90.0	69.5	84.3	84.4	73.5	64.2	57.2	41.8	66.3	60.7	78.0	70.9	70.3
w. RIL (Dr.GRPO + GAIL)	86.7	95.4	74.5	86.8	87.2	77.3	66.1	61.8	48.2	71.1	68.5	80.7	73.0	74.2
Qwen2.5-VL-3B	81.6	84.0	61.2	79.1	78.1	61.8	57.6	51.2	31.6	55.9	47.6	74.0	67.6	65.4
w. RL (GRPO)	81.9	87.8	61.8	80.3	79.7	62.8	55.4	52.0	33.7	57.7	52.6	75.7	68.6	66.4
w. RL (Dr.GRPO)	82.4	89.7	62.5	81.8	81.2	63.8	55.9	52.5	34.3	59.1	53.7	76.5	69.2	67.0
w. RIL (Dr.GRPO + GAIL)	83.0	95.4	65.2	84.8	84.2	67.4	57.6	53.7	36.8	61.2	55.5	78.7	70.3	67.7
InternVL3-8B	85.2	86.6	71.6	83.4	82.2	78.5	66.3	62.7	41.3	68.2	55.5	77.1	69.7	70.8
w. RL (GRPO)	85.9	89.6	72.3	84.6	84.2	78.7	66.4	63.8	42.5	70.0	57.1	77.9	70.6	71.3
w. RL (Dr.GRPO)	86.3	91.2	72.9	85.2	85.3	79.0	66.5	64.9	42.8	70.7	57.5	78.1	71.1	72.0
w. RIL (Dr.GRPO + GAIL)	87.4	95.5	74.1	88.7	89.3	78.4	66.5	66.8	44.8	75.4	60.1	80.6	73.0	73.3
InternVL3-2B	78.7	80.2	57.0	81.1	78.4	62.2	53.9	48.6	24.9	60.7	47.0	75.0	64.6	64.3
w. RL (GRPO)	79.1	86.9	57.8	82.4	80.0	61.8	53.6	49.6	25.7	61.5	47.3	75.8	64.9	64.6
w. RL (Dr.GRPO)	79.8	90.3	58.5	83.2	81.3	62.1	53.8	50.5	26.4	62.4	47.9	76.5	65.4	65.2
w. RIL (Dr.GRPO + GAIL)	81.1	93.6	63.0	85.6	83.2	61.5	54.2	52.9	27.5	63.2	48.3	78.2	66.7	66.8
InternVL3-1B	69.4	75.3	45.8	72.6	67.9	58.7	47.5	43.4	17.5	51.5	42.9	71.2	58.2	58.2
w. RL (GRPO)	69.6	83.0	46.7	73.8	69.4	57.7	47.0	42.8	17.2	52.5	42.7	71.5	58.5	58.5
w. RL (Dr.GRPO)	69.9	86.7	47.6	74.9	70.5	57.5	47.0	42.8	17.3	53.4	42.9	71.9	59.4	59.4
w. RIL (Dr.GRPO + GAIL)	72.0	93.6	51.6	79.6	75.7	56.9	47.4	43.1	18.3	57.4	44.1	74.2	61.4	61.4

student VLMs are correct, and (4) utilizing the generated text responses from teacher VLMs when GRPO [43] updates student VLMs, serving as a direct guidance to reach the correct text responses. These approach not only stabilizes RIL training but also provides a mechanism for student VLMs to potentially exceed the performance of teacher VLMs. Moreover, RIL do not require an explicit think-answer process, thereby maintaining the same inference, and it is independent of any particular image embedding methods or language tokenizers used in the student or the teacher VLMs.

3 Unified Reinforcement and Imitation Learning for VLMs

3.1 Core Model Components for RIL

Our RIL framework utilizes several key model components. For the student models—which are the VLMs we aim to enhance—we employ recently released architectures such as Qwen2.5-VL (3B and 7B variants) [1] and InternVL3 (1B, 2B, and 8B variants) [18]. Starting from these pre-trained checkpoints, RIL directly updates their parameters with the goal of developing efficient, highperforming VLMs capable of rivaling leading closed-source systems like GPT-40 [50], Gemini [51], and Claude-3.5 Sonnet [52]. For the teacher models, which provide the target behavior, we select powerful large teacher VLMs recognized for their strong performance on benchmarks like VLM Leaderboard [68], specifically Qwen2.5-VL-72B [1] and InternVL3-78B [18]. The student VLMs are trained to mimic the text generation patterns and response styles of these teacher models. Central to the imitation learning aspect is a discriminator, designed to assess the similarity between studentgenerated responses and those from the teacher VLMs. To promote training stability and mitigate the balance problem common in adversarial setups [47, 48]—where one component might overpower the other—the discriminator's architecture and initial parameters deliberately mirror those of the student VLMs. Lastly, to evaluate the factual correctness of the generated text, we utilize an LLMas-a-Judge [9], specifically Qwen2.5-32B [13], following the methodology of Zheng et al. [9]. This model is not trained but is used solely to determine whether the predicted response accurately reflects the meaning of the ground truth.

The subsequent sections will first elaborate on the discriminator's architecture and pre-training (Section 3.2), followed by a detailed exposition of the RIL algorithm and its reward design (Section 3.3).

3.2 Discriminator Architecture and Pre-training

With the student (small) VLMs serving as the generator in our framework, the effective pre-training of the discriminator is crucial. This initial training phase equips the discriminator (D_{ϕ}) with the ability to reliably differentiate between text responses originating from student VLMs versus teacher

Table 2: Showing the effectiveness of employing multiple large VLMs more than a single VLM, under the numerous evaluation benchmarks equally used in Table 1.

VLMs	AI2D	ChartQA	MathVista	MMB	MMB ^{CN}	MM-Vet	MM-Vet-v2	2 MMMU	MMMU-Pro	MMStar	BLINK	SEED	SEED2+	RWQA
Qwen2.5-VL-7B	83.9	87.3	67.8	83.5	83.4	71.8	63.7	55.0	38.3	63.9	56.4	77.0	70.4	68.5
w. RIL (Qwen2.5-VL-72B)	86.7	95.4	74.5	86.8	87.2	77.3	66.1	61.8	48.2	71.1	68.5	80.7	73.0	74.2
w. RIL (InternVL3-78B)	86.8	95.5	74.6	86.7	87.1	75.8	66.0	60.9	47.1	71.1	68.1	80.8	72.7	75.4
w. RIL (Both)	86.1	95.6	79.7	86.3	86.5	80.4	71.1	65.7	48.5	71.1	70.0	80.5	72.8	72.8
Qwen2.5-VL-3B	81.6	84.0	61.2	79.1	78.1	61.8	57.6	51.2	31.6	55.9	47.6	74.0	67.6	65.4
w. RIL (Qwen2.5-VL-72B)	83.0	95.4	65.2	84.8	84.2	67.4	57.6	53.7	36.8	61.2	55.5	78.7	70.3	67.7
w. RIL (InternVL3-78B)	83.3	95.4	65.1	85.0	84.7	66.6	56.3	54.9	37.7	61.1	55.3	78.8	70.2	68.2
w. RIL (Both)	83.9	95.7	71.7	85.2	84.5	74.7	62.8	60.7	43.7	65.2	60.8	78.7	71.4	73.2
InternVL3-8B	85.2	86.6	71.6	83.4	82.2	78.5	66.3	62.7	41.3	68.2	55.5	77.1	69.7	70.8
w. RIL (Qwen2.5-VL-72B)	87.5	95.2	74.3	88.8	89.8	77.6	65.9	67.3	45.7	75.2	59.1	80.4	72.9	72.5
w. RIL (InternVL3-78B)	87.4	95.5	74.1	88.7	89.3	78.4	66.5	66.8	44.8	75.4	60.1	80.6	73.0	73.3
w. RIL (Both)	87.3	95.3	77.8	88.1	89.6	80.1	67.6	68.6	47.8	74.8	62.7	80.5	73.8	73.7
InternVL3-2B	78.7	80.2	57.0	81.1	78.4	62.2	53.9	48.6	24.9	60.7	47.0	75.0	64.6	64.3
w. RIL (Qwen2.5-VL-72B)	81.1	93.6	61.9	85.3	83.0	62.6	52.8	52.6	26.6	63.6	48.8	78.0	67.1	66.3
w. RIL (InternVL3-78B)	81.1	93.6	63.0	85.6	83.2	61.5	54.2	52.9	27.5	63.2	48.3	78.2	66.7	66.8
w. RIL (Both)	80.7	94.0	67.4	84.7	82.2	70.7	58.4	56.6	32.3	63.9	51.5	77.8	66.8	69.3
InternVL3-1B	69.4	75.3	45.8	72.6	67.9	58.7	47.5	43.4	17.5	51.5	42.9	71.2	58.2	58.2
w. RIL (Qwen2.5-VL-72B)	72.2	93.6	50.9	79.9	75.8	56.1	47.3	43.6	18.0	58.0	43.6	74.6	60.7	61.0
w. RIL (InternVL3-78B)	72.0	93.6	51.6	79.6	75.7	56.9	47.4	43.1	18.3	57.4	44.1	74.2	61.4	61.4
w. RIL (Both)	73.0	94.1	55.5	79.1	75.9	62.9	50.7	49.7	19.0	60.5	46.9	74.7	61.9	63.7

(large) VLMs. Without adequate pre-training, D_{ϕ} would yield unreliable scores, undermining its utility and potentially degrading overall model performance. To prepare for pre-training, we first collect a dataset of responses. For each given question q, N text responses are generated from both the student VLMs, denoted as $\{\mathbf{o}_i^{(s)}\}_{i=1}^N$, and the teacher VLMs, denoted as $\{\mathbf{o}_i^{(t)}\}_{i=1}^N$. Using this collection, the discriminator D_{ϕ} , parameterized by ϕ , is trained to maximize the following objective:

$$\max_{\phi} \mathcal{L}(\phi) = \frac{1}{N} \sum_{i=1}^{N} \left\{ \log D_{\phi}(q, o_i^{(s)}) + \log(1 - D_{\phi}(q, o_i^{(t)})) \right\}. \tag{1}$$

According to this objective, D_{ϕ} learns to output a score approaching 'one value' for responses from student VLMs and 'zero value' for those from teacher VLMs. For practical implementation, given that the discriminator's backbone is language-based (mirroring the student VLMs), the generated responses o are formatted using a specific prompt (detailed in Fig. 3) to ensure proper textual understanding. To elicit the scalar discrimination score, the standard language head of the VLM (typically mapping to vocabulary logits, $(\mathbb{R}^{d \times v})$ is replaced with a linear discriminator head $(\mathbb{R}^{d \times 1})$. The input to this head is the representation of the final sequence token from the last layer of the backbone model, from which the score is directly computed with the sigmoid function.

3.3 Mimicking Large Teacher VLMs with Similarity and Answer Rewards

The RIL process commences after an initial supervised finetuning (SFT) phase for the student (small) VLMs π_{θ} using a comprehensive visual instruction tuning dataset (see Appendix C). This SFT stage acts as a crucial warm-up, acclimating π_{θ} to the target data distribution. Once the student VLMs and the pre-trained discriminator D_{ϕ} (from Section 3.2) are prepared, the RIL loop begins. For each question q in a training batch, we generate G text responses from the current student VLMs, $\{o_i^{(s)}\}_{i=1}^G$, and retrieve G responses corresponding the question q from the teacher (large) VLMs, $\{o_i^{(t)}\}_{i=1}^G$. We note that pre-generating and caching the responses from teacher VLMs (e.g., during the discriminator's pre-training) can significantly reduce RIL training time. In the end, this yields a combined set of 2G responses, $\{o_i\}_{i=1}^{2G}$, for each question. Within each RIL iteration, we first update the discriminator D_{ϕ} using Eq. (1) on these 2G responses. This step ensures that D_{ϕ} maintains its ability to distinguish between student and teacher outputs as the student VLMs evolve. Next, the student VLMs π_{ϕ} are updated using an objective derived from Dr.GRPO [8], an advanced variant of GRPO that provides unbiased advantage estimates $\{\hat{A}_i\}_{i=1}^{2G}$. The objective function is:

$$\max_{\theta} \mathcal{L}(\theta) = \frac{1}{2G} \sum_{i=1}^{2G} \left\{ \min \left[r_i(\theta) \hat{A}_i, \operatorname{clip} \left(r_i(\theta), 1 - \epsilon, 1 + \epsilon \right) \hat{A}_i \right] - \beta \mathcal{D}_{KL} \left(\pi_{\theta} \mid \pi_{\text{ref}} \right) \right\},$$
where
$$r_i(\theta) = \frac{\pi_{\theta}(o_i \mid q)}{\pi_{\theta_{\text{old}}}(o_i \mid q)}, \quad \hat{A}_i = R(q, o_i) - \frac{1}{2G} \sum_{j=1}^{2G} R(q, o_j).$$
(2)

Here, $r_i(\theta)$ is the probability ratio between the current policy π_{θ} and the policy before the update π_{old} , π_{ref} is typically the initial SFT model, ϵ is a clipping hyperparameter, and β controls the KL-

Table 3: Comparing **RIL**-applied VLMs with standard or smaller model size open-source VLMs.

VLMs	AI2D	ChartQA	MathVista	${\rm MMB}$	MMB^{CN}	MM-Vet	$\mathbf{M}\mathbf{M}\mathbf{M}\mathbf{U}$	MMMU-Pro	MMStar	BLINK	SEED	SEED2+	RWQA
LLaVA-OneVision-7B [28]	81.4	80.0	63.2	80.8	-	57.5	48.8	24.1	61.9	53.0	76.7	65.4	69.9
InternVL2-8B [69]	83.8	83.3	58.3	81.7	81.2	54.2	49.3	29.0	61.5	50.9	75.4	67.3	64.2
MiniCPM-V2.5-8B [70]	78.4	-	54.3	77.2	74.2	52.8	45.8	-	51.8	-	72.3	61.4	63.5
MiniCPM-V2.6-8B [70]	82.1	-	60.6	-	-	60.0	49.8	27.2	57.5	55.2	74.0	65.7	65.0
MiniCPM-o2.6-8B [70]	86.1	86.9	73.3	-	-	67.2	50.9	-	63.3	53.9	75.5	67.9	68.0
Ovis2-8B [71]	86.6	-	71.8	-	-	65.1	57.4	-	64.6	54.3	77.2	70.4	72.5
Ovis2-16B [71]	86.3	-	73.7	-	-	68.4	60.7	-	67.2	59.0	77.7	72.1	74.1
Qwen2-VL-7B [16]	77.5	83.0	58.2	83.0	80.5	62.0	54.1	30.5	60.7	53.8	76.0	68.6	68.5
InternVL2.5-8B [17]	84.8	84.8	64.4	84.6	82.6	62.8	56.0	34.3	63.2	54.8	76.8	69.7	70.1
Qwen2.5-VL-7B [1]	83.9	87.3	67.8	83.5	83.4	71.8	55.0	38.3	63.9	56.4	77.0	70.4	68.5
InternVL3-8B [18]	85.2	86.6	71.6	83.4	82.2	78.5	62.7	41.3	68.2	55.5	77.1	69.7	70.8
Qwen2.5-VL-RIL-7B (Both)	86.1	95.6	79.7	86.3	86.5	80.4	65.7	48.5	71.1	70.0	80.5	72.8	72.8
InternVL3-RIL-8B (Both)	87.3	95.3	77.8	88.1	89.6	80.1	68.6	47.8	74.8	62.7	80.5	73.8	73.7
Phi-3.5-Vision-4B [72]	77.8	81.8	43.9	76.0	66.1	43.2	43.0	19.7	47.5	58.3	69.7	62.2	53.6
Phi-4-Multimodal-5.6B [73]	83.0	81.4	65.8	86.7	-	51.9	56.0	38.5	58.9	42.1	73.2	68.5	64.1
Ovis2-4B [71]	85.7	-	69.6	-	-	65.5	49.0	-	61.9	53.0	76.2	69.3	71.1
InternVL2-4B [69]	78.9	81.5	58.6	78.6	73.9	51.0	34.3	32.7	53.9	46.1	73.2	63.9	60.7
InternVL2.5-4B [69]	81.4	84.0	60.5	81.1	79.3	60.6	52.3	32.7	58.7	50.8	74.8	66.9	64.3
Qwen2.5-VL-3B [1]	81.6	84.0	61.2	79.1	78.1	61.8	51.2	31.6	55.9	47.6	74.0	67.6	65.4
Qwen2.5-VL-RIL-3B (Both)	83.9	95.7	71.7	85.2	84.5	74.7	60.7	43.7	65.2	60.8	78.7	71.4	73.2
InternVL2-2B [69]	74.1	76.2	46.3	73.2	70.9	39.5	34.3	18.2	49.8	43.8	70.9	60.0	57.3
Qwen2-VL-2B [16]	60.2	73.5	43.0	74.9	73.5	49.5	41.1	21.2	47.5	45.2	72.4	61.2	62.6
Aquila-VL-2B [17]	75.0	76.5	59.0	-	-	43.8	47.4	-	54.9	34.1	73.9	63.0	65.0
Ovis2-2B [71]	82.7	-	64.1	-	-	58.3	45.6	-	56.7	47.9	74.4	67.4	66.0
InternVL2.5-2B [17]	74.9	79.2	51.3	74.7	71.9	60.8	43.6	23.7	54.3	44.0	73.2	60.0	60.1
InternVL3-2B [18]	78.7	80.2	57.0	81.1	78.4	62.2	48.6	24.9	60.7	47.0	75.0	64.6	64.3
InternVL3-RIL-2B (Both)	80.7	94.0	67.4	84.7	82.2	70.7	56.6	32.3	63.9	51.5	77.8	66.8	69.3
LLaVA-OneVision-0.5B [28]	57.1	61.4	34.8	61.6	55.5	32.2	31.4	-	37.7	52.1	66.6	45.7	55.6
Ovis2-1B [71]	76.4	-	59.4	-	-	50.0	36.1	-	52.1	44.0	71.7	61.4	63.9
InternVL2-1B [69]	64.1	72.9	37.7	65.4	60.7	32.7	36.7	14.8	45.6	38.6	65.2	54.3	50.3
InternVL2.5-1B [17]	69.3	75.9	43.2	70.7	66.3	48.8	40.9	18.8	51.3	42.0	70.4	59.0	57.5
InternVL3-1B [18]	69.4	75.3	45.8	72.6	67.9	58.7	43.4	17.5	51.5	42.9	71.2	58.2	58.2
InternVL3-RIL-1B (Both)	73.0	94.1	55.5	79.1	75.9	62.9	49.7	19.0	60.5	46.9	74.7	61.9	63.7

divergence penalty. The total reward $R(q, o_i)$ for each response o_i is a crucial composite signal, comprising two binary components:

- 1. **Similarity Reward**: This indicates if the response o_i resembles those from teacher VLMs, denoted by $\mathbb{1}(D_\phi(q,o_i)<0.5)$. A discriminator score below 0.5 suggests higher similarity to teacher responses, consistent with D_ϕ being trained to output 'zero value' for teacher samples (Eq. 1).
- 2. **Answer Reward**: It assesses the factual correctness of o_i against a ground truth answer a, determined by LLM-as-a-Judge (q, a, o_i) . This evaluation uses the prompt detailed in Fig. 4.

These two rewards are summed to form the overall reward $R(q,o_i)$. The inclusion of the answer reward is essential, as the discriminator alone primarily captures stylistic similarity and does not guarantee factual accuracy, the absence of which could lead to significant performance issues. This dual-reward strategy ensures that mimicking large teacher VLMs involves not just replicating surface-level text patterns but also encompasses the deeper 'verbalization effect' [45], which characterizes how proficient VLMs articulate correct and contextually appropriate answers. The comprehensive RIL procedure is detailed in Algorithm 2.

4 Experiments

4.1 Dissecting RIL

To understand the distinct contributions of its components, we first analyze RIL by comparing it against reinforcement learning (RL) baselines that utilize only answer rewards. Specifically, we evaluate RL agents trained solely with GRPO [43] or its advanced variant, Dr.GRPO [8], without the discriminator or the associated similarity rewards. As shown in Table 1, RIL, which integrates reinforcement and imitation learning elements based on Dr.GRPO and GAIL [44], demonstrably outperforms these RL-only approaches, confirming the efficacy of our unified training algorithm.

A key finding is the substantial benefit of employing multiple large teacher VLMs. As detailed in Table 2, using responses from a combination of large models (specifically Qwen2.5-VL-72B [1] and InternVL3-78B [18]) yields significantly better vision-language performance than relying on any single teacher. This advantage arises from the richer diversity of textual responses, which strengthens the discriminator's training. Furthermore, incorporating this diverse set of high-quality responses from teacher VLMs directly into the GRPO optimization step provides student VLMs with clearer and

T 1 1 4 C	1 3 77 3 6 1.1 1	1 1 1 1 1 1 1 1
Table 4: Comparing RII -appli	ed VIIVIs with large size o	pen-source and closed-source VLMs.
Table 7. Companing IXID-appi	ca v Livis with large size o	pen-source and crosed-source villyis.

VLMs	AI2D	ChartQA	MathVista	a MMB	MM-Vet	MM-Vet-v2	2 MMMU	MMMU-Pro	MMStar	BLINK	SEED	SEED2+	RWQA
NVLM-72B [54]	85.2	86.0	66.6	-	58.9	-	59.7	-	63.7	48.0	75.5	68.4	69.9
LLaVA-OV-72B [28]	85.6	83.7	67.5	85.8	60.6	-	56.8	31.0	65.8	55.4	77.5	-	71.9
Molmo-72B [53]	83.4	87.3	58.6	-	61.1	-	54.1	-	63.3	49.7	74.6	67.6	73.7
Qwen2-VL-72B [16]	88.1	88.3	70.5	86.5	74.0	68.7	64.5	46.2	68.3	60.5	77.9	72.3	77.8
Qwen2.5-VL-72B [1]	88.7	89.5	74.2	88.6	76.9	76.2	68.2	46.2	70.8	64.4	79.5	73.0	75.7
InternVL2-76B [69]	87.6	88.4	65.5	86.5	65.7	68.4	62.7	40.0	67.4	56.8	77.6	69.7	72.2
InternVL2.5-78B [28]	89.1	88.3	72.3	88.3	72.3	65.5	70.1	48.6	69.5	63.8	77.0	71.3	78.7
InternVL3-78B [18]	89.7	89.7	79.0	89.0	81.3	70.0	72.2	47.3	72.5	66.3	78.7	71.9	78.0
Claude-3.5-Sonnet [52]	81.2	90.8	67.7	82.6	70.1	71.8	68.3	51.5	65.1	60.1	61.7	71.7	65.8
Claude-3.7-Sonnet [52]	82.5	-	66.8	-	70.0	-	71.0	-	65.1	56.6	74.3	67.6	55.4
Gemini-1.5-Pro [51]	79.1	87.2	63.9	73.9	64.0	66.9	62.2	46.9	59.1	59.1	76.0	70.8	67.5
Gemini-2.0-Flash [51]	83.1	-	70.4	90.0	73.6	-	69.9	54.4	69.4	64.0	-	-	72.3
GPT-4o (24.05.13) [74]	84.6	85.7	63.8	83.4	69.1	71.0	69.1	51.9	64.7	68.0	77.1	72.0	75.4
GPT-4.1 (25.04.14) [74]	85.9	-	70.4	-	78.8	-	74.0	-	69.8	68.5	78.0	73.1	78.7
Qwen2.5-VL-RIL-7B (Both)	86.1	95.6	79.7	86.3	80.4	71.1	65.7	48.5	71.1	70.0	80.5	72.8	72.8
InternVL3-RIL-8B (Both)	87.3	95.3	77.8	88.1	80.1	67.6	68.6	47.8	74.8	62.7	80.5	73.8	73.7




Figure 5: Illustrating training dynamics of small VLMs during RIL, where (Left) showing the evolution of similarity rewards over training iterations and (Mid) accuracy rewards obtained using LLM-as-a-Judge [9], ensuring that generated responses are both contextually appropriate and factually correct. (Right) displaying the overall average performance of evaluation benchmarks used in Table 1.

more varied exemplars of correct answer generation. To be simplified in the manner of an algorithm, given a response 'T1' from one teacher VLM and 'T2' from another teacher VLM, and a response 'S' from the student VLM, then the discriminator is trained to output zero when 'S' is provided and one when either 'T1' or 'T2' is provided. Additionally, the rewards and their advantages for 'S', 'T1', and 'T2' are computed based on similarity and answer quality in order to train the student VLM by using RIL's objective loss.

Beyond these component analyses, our RIL-trained student VLMs exhibit highly competitive performance against a range of recent open-source and closed-source models. Comparisons across various model sizes (detailed in Table 3 for smaller models and Table 4 for larger ones) and diverse vision-language benchmarks (illustrated in Figures 1 and 2) show that RIL not only substantially narrows the performance gap to state-of-the-art VLMs but, in several instances, surpasses them.

Finally, the training dynamics of RIL, depicted in Figure 5, offer further insight. Over the course of training, RIL consistently increases both the similarity rewards (from the discriminator) and the answer rewards (from LLM-as-a-Judge [9]). It indicates that the student VLMs are progressively generating responses that are both more stylistically aligned with teacher models and factually correct, with these improving reward signals directly translating into significant overall performance gains.

For a comprehensive account of the specific training procedures, hyperparameter settings, evaluation setups, and dataset details that ensure the reproducibility of these, please refer to Appendix C. Using 256 NVIDIA A100 GPUs, pre-training the discriminator on 1.2M samples (40K [number of samples] \times 16 [generated responses] \times 2 [for both teacher and student]) takes approximately 1 to 3 days. The SFT step on the 4M-sample SFT dataset takes around 3 to 5 days. Conducting the RIL loop for the sampled 40K data requires an additional 3 to 5 days using 8 NVIDIA A100 GPUs.

4.2 Impacts of Hyperparameters, Techniques, and Distillation

We systematically investigate the sources of RIL's effectiveness by analyzing key training hyperparameters, component design choices, and its interplay with other learning techniques. Note that, for the optimization objective of student VLMs in Eq. (2), the clipping hyperparameter ϵ is consistently set to 0.2. Our first analyses focus on core optimization settings. We examine the number of update iterations, μ , for the discriminator and student VLMs within each RIL cycle (related to Algorithm 2).

Table 5: Analyzing the impacts of training hyperparameters, techniques for stabilization, and the effect of distillation. Note that, these tables (a)-(c) set a backbone model as Qwen2.5-VL-7B [1] and apply RIL to this model in order to validate it by controlling multiple factors .

(a) l														1				
(a) Discriminator & small VLMs Iteration (μ)							(b)	KL-	livergeno	e Pen	alty Co	efficie	(c) Importance of VLM Modules					
μ	AI2D	MathVista	MMB	MM-Vet	MMMU	BLINK	β	AI2D	MathVista	MMB	MM-Vet	MMMU	BLINK	Module	MathVista	MMB	MM-Vet	MMMU
1	86.1	79.7	86.3	80.4	65.7	70.0	0.00	77.1	59.9	73.2	60.2	42.9	45.0	Full Training	79.7	86.3	80.4	65.7
2	86.1	79.0	86.3	80.4	65.3	69.2	0.01	86.1	79.7	86.3	80.4	65.7	70.0	w.o. Word-Em	bed 75.4	85.5	75.1	62.2
4	85.5	78.8	85.5	79.5	65.0	68.4	0.05	85.7	77.3	85.7	78.5	63.5	68.3	w.o. Self-Attn		84.3	75.2	61.3
6	85.1	77.5	85.2	78.0	63.8	68.0	0.10	85.2	74.9	85.2	76.6	63.0	65.8	w.o. FFN-Gate		86.0	78.8	63.5
8	84.5	76.2	84.7	77.4	62.5	66.7	0.20	84.8	74.4	84.5	75.8	62.5	64.2	w.o. FFN-Up	77.3	86.0	79.0 79.6	64.1 64.8
10	84.3	75.0	84.5	77.2	61.9	65.5	0.50	84.3	74.0	84.8	75.9	61.3	62.0	w.o. FFN-Dow w.o. Layer-No		86.1 86.3	79.6 80.4	65.7
12		73.8	84.4	76.9	61.0	64.2			73.9	84.3	74.8	60.4	61.7	w.o. Lang-Hea		85.2	73.9	62.0
	(d) Importance of discriminator D_{ϕ}																	
	(d) I1	mportan	ce of c	liscrimi	inator <i>L</i>	O_{ϕ}	(e) E	ffect	of Answ	er Re	wards ((AR) aı	nd SFT	(1	f) Effect of I	Distill	ation	
VL	` /				inator <i>I</i>	Ψ	(e) E		of Answ	er Re				VLMs (1	f) Effect of I	Distill MMB	ation MM-Vet	MMMU
VL	` /		MathVis	ta MMB	MM-Vet	MMMU					ista MMB			VLMs	RIL MathVista X 61.3	MMB 84.4	MM-Vet 65.1	57.4
_	Ms	D_{ϕ}	MathVis 69.5	ta MMB 84.3	MM-Vet 73.5	MMMU 57.2	VLMs		AR	MathV 65.2	ista MMB ! 82.1 8 82.5	MM-Vet 69.9 70.3	MMMU 53.6 53.8	VLMs	RIL MathVista x 61.3 √ 70.8	MMB 84.4 85.3	MM-Vet 65.1 73.3	57.4 66.1
_	` /	D_{ϕ}	MathVis 69.5 75.3	84.3 85.5	MM-Vet 73.5 75.1	MMMU 57.2 61.2	VLMs		AR	MathV 65.2 65.3 79.7	ista MMB 82.1 82.5 86.3	MM-Vet 69.9 70.3 80.4	MMMU 53.6 53.8 65.7	VLMs	RIL MathVista X 61.3 70.8 X 61.5	MMB 84.4 85.3 84.8	MM-Vet 65.1 73.3 65.3	57.4 66.1 57.9
_	Ms	<i>D</i> _φ /L-7B Δ /	MathVis 69.5 75.3 79.7	84.3 85.5 86.3	73.5 75.1 80.4	MMMU 57.2 61.2 65.7	VLMs		AR	MathV 65.2 65.3 79.7 Γ 73.5	ista MMB 2 82.1 3 82.5 7 86.3 5 84.1	MM-Vet 69.9 70.3 80.4 76.4	53.6 53.8 65.7 60.7	VLMs MiniLLM-7B DistilLLM-7B	RIL MathVista X 61.3 V 70.8 X 61.5 V 71.0	MMB 84.4 85.3 84.8 85.4	MM-Vet 65.1 73.3 65.3 73.9	57.4 66.1 57.9 67.2
Qw	Ms ven2.5-V	D _φ /L-7B Δ ✓	MathVis 69.5 75.3 79.7 72.9	84.3 85.5 86.3 85.2	73.5 75.1 80.4 79.0	MMMU 57.2 61.2 65.7 64.9	VLMs Qweni	2.5-VL-7	AR	MathV 65.2 65.3 79.7 73.5 69.3	ista MMB 2 82.1 3 82.5 7 86.3 6 84.1 8 81.4	MM-Vet 69.9 70.3 80.4 76.4 75.4	53.6 53.8 65.7 60.7 59.3	VLMs MiniLLM-7B	RIL MathVista X 61.3 70.8 X 61.5	MMB 84.4 85.3 84.8	MM-Vet 65.1 73.3 65.3	57.4 66.1 57.9
Qw	Ms	D _φ /L-7B Δ ✓	MathVis 69.5 75.3 79.7	84.3 85.5 86.3	73.5 75.1 80.4	MMMU 57.2 61.2 65.7	VLMs Qweni		AR	MathV 65.2 65.3 79.7 Γ 73.5	ista MMB 2 82.1 3 82.5 7 86.3 6 84.1 6 81.4 0 82.1	MM-Vet 69.9 70.3 80.4 76.4	53.6 53.8 65.7 60.7	VLMs MiniLLM-7B DistilLLM-7B	RIL MathVista X 61.3 Y 70.8 X 61.5 Y 71.0 X 61.8	MMB 84.4 85.3 84.8 85.4 85.0	MM-Vet 65.1 73.3 65.3 73.9 65.9	57.4 66.1 57.9 67.2 58.2

As shown in Table 5(a), a single update iteration ($\mu=1$) for both is sufficient for strong performance; more iterations, especially for student VLM updates, can risk overfitting. Subsequently, we explore the impact of the KL-divergence penalty coefficient β (Eq. 2) in Table 5(b). These results confirm that while constraining policy updates is crucial for stability, an overly restrictive penalty (large β) can hinder performance.

Next, we assess decisions related to model component updates. Table 5(c) details the importance of updating different parameter groups within the student VLMs, revealing that training the self-attention layers, word embeddings, and the language head yields more significant gains than updating feed-forward networks (FFN) or layer normalization parameters. We also investigate the necessity of continuously updating the discriminator during RIL training (Table 5(d)). Interestingly, Table 5(d) shows that using a fixed (not updated: \triangle), pre-trained discriminator can outperform RL-only baselines (denoted by \checkmark), but overall RIL efficacy still critically depends on leveraging a continuously trained discriminator (denoted by \checkmark) during training RIL.

In addition, we evaluate RIL's interaction with evaluation methodologies of the predicted answer and its synergy with other training paradigms. Table 5(e) underscores the importance of LLM-as-a-Judge [9] for answer rewards; replacing it with conventional answer parsing (denoted \triangle) significantly degrades performance, particularly on general visual questions (e.g., "How to cook this dish?" from an image) where fixed parsing rules are inadequate. This highlights the need for flexible, LLM-based evaluation. The same table also affirms that while initial SFT is beneficial for capturing the training data distribution, RIL subsequently delivers substantially larger improvements.

On Table 5(f), we examine RIL's effectiveness on student VLMs that are already trained on distillation methods like MiniLLM [75], DistilLLM [76], LLaVA-KD [77], and VLsI [45] (on same backbones like Qwen2-VL [16] and training datasets C). As shown in Table 5(f), RIL proves particularly potent for these models compared to non-distillation counterparts, likely because prior feature alignment through distillation primes them for RIL's alignment mechanisms.

Lastly, we demonstrate that the binary similarity reward introduced in Sec. 3.3 enables more efficient and stable training of RIL. As shown in Fig. 6, the binary reward consistently outperforms continuous and multi-level discretizations, suggesting that binary feedback provides clearer and more reliable learning signals. This improvement arises because models often struggle to interpret subtle differences in continuous rewards (e.g., why a score of 0.21 should be preferred over 0.20).

4.3 Discussion and Limitation

RIL leverages the diversity of responses generated by large VLMs, which tend to produce more varied and higher-quality answers beyond the ground truth answers used in SFT training datasets. As illustrated in Fig. 7, increasing the number of teacher-generated responses consistently improves the student model's performance, demonstrating that richer supervision from diverse teacher outputs enhances the student VLM's overall generalization ability. In summary, RIL exploits these diverse, high-quality teacher responses to guide smaller VLMs toward learning more flexible responses beyond fixed ground truth answers.

Furthermore, RIL offers distinct advantages over traditional knowledge distillation techniques [78]. Unlike conventional distillation, which typically requires student VLMs to replicate the high-




Figure 6: Comparison of RIL performance using different discretization levels of the similarity reward on four benchmarks: MathVista [3], MMB [20], MM-Vet [4], and MMMU [5]) with Qwen2.5-VL-7B [1] and InternVL3-8B [18].



Figure 7: Performance of RIL with varying numbers of generated responses from a teacher VLM across four benchmarks and their average. They are evaluated on Qwen2.5-VL-7B [1].

dimensional logit features of teacher models, RIL operates exclusively on generated text responses in natural language. This fundamental difference liberates RIL from restrictive feature-level constraints, such as the need for identical image embedding strategies or perfectly aligned language tokenizers (including vocabulary, index order, and sequence length) between student and teacher VLMs. Common distillation methods [75–77, 45] often rely on KL divergence—a static, non-trainable metric devoid of contextual understanding—to measure feature similarity, but RIL employs a trainable discriminator which assesses natural language responses directly, learning to discern nuanced similarities and stylistic differences between student and teacher outputs. The adaptability of the trainable discriminator in capturing contextual relationships makes RIL a particularly flexible and powerful approach for lightweight VLM and enhancement.

While RIL demonstrates these significant strengths, its current implementation has been focused primarily on the post-instruction tuning alignment phase. A promising avenue for future research, therefore, involves extending the utility of our discriminator to the initial visual instruction tuning stage itself. We hypothesize that the discriminator's ability to learn from natural language signals and overcome conventional architectural constraints could also offer significant benefits in this earlier phase of VLM training, potentially leading to more efficient instruction following from the outset.

5 Conclusion

In this work, we have introduced Unified Reinforcement and Imitation Learning (RIL), a novel framework enabling smaller VLMs to emulate the sophisticated text-generation capabilities of substantially larger counterparts, often matching or even exceeding their performance. RIL directly addresses the pressing need for high-performing, lightweight VLMs deployable in resource-constrained settings. By leveraging a dual reward system—combining similarity scores from a discriminator with answer accuracy assessments from an LLM-as-a-Judge—RIL effectively guides student VLMs towards superior alignment and overall efficacy. Key findings from our research include the significant performance boost achieved by using multiple large teacher VLMs, which provides a richer diversity of training signals, and the notable success of RIL in enhancing distillation-based VLMs. Comprehensive experiments across a wide array of vision-language benchmarks confirm that RIL substantially narrows, and in several cases surpasses, the performance gap to state-of-the-art openand closed-source VLMs.

References

- [1] S. Bai, K. Chen, X. Liu, J. Wang, W. Ge, S. Song, K. Dang, P. Wang, S. Wang, J. Tang, et al., "Qwen2. 5-vl technical report," arXiv preprint arXiv:2502.13923, 2025.
- [2] A. Kembhavi, M. Salvato, E. Kolve, M. Seo, H. Hajishirzi, and A. Farhadi, "A diagram is worth a dozen images," in *Computer Vision–ECCV 2016: 14th European Conference, Amsterdam, The Netherlands, October 11–14*, 2016, *Proceedings, Part IV 14*, pp. 235–251, Springer, 2016.
- [3] P. Lu, H. Bansal, T. Xia, J. Liu, C. Li, H. Hajishirzi, H. Cheng, K.-W. Chang, M. Galley, and J. Gao, "Mathvista: Evaluating mathematical reasoning of foundation models in visual contexts," *arXiv preprint arXiv:2310.02255*, 2023.
- [4] W. Yu, Z. Yang, L. Li, J. Wang, K. Lin, Z. Liu, X. Wang, and L. Wang, "Mm-vet: Evaluating large multimodal models for integrated capabilities," *arXiv preprint arXiv:2308.02490*, 2023.
- [5] X. Yue, Y. Ni, K. Zhang, T. Zheng, R. Liu, G. Zhang, S. Stevens, D. Jiang, W. Ren, Y. Sun, *et al.*, "Mmmu: A massive multi-discipline multimodal understanding and reasoning benchmark for expert agi," *arXiv preprint arXiv:2311.16502*, 2023.
- [6] X. Fu, Y. Hu, B. Li, Y. Feng, H. Wang, X. Lin, D. Roth, N. A. Smith, W.-C. Ma, and R. Krishna, "Blink: Multimodal large language models can see but not perceive," arXiv preprint arXiv:2404.12390, 2024.
- [7] D. Guo, D. Yang, H. Zhang, J. Song, R. Zhang, R. Xu, Q. Zhu, S. Ma, P. Wang, X. Bi, *et al.*, "Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning," *arXiv* preprint arXiv:2501.12948, 2025.
- [8] Z. Liu, C. Chen, W. Li, P. Qi, T. Pang, C. Du, W. S. Lee, and M. Lin, "Understanding r1-zero-like training: A critical perspective," *arXiv preprint arXiv:2503.20783*, 2025.
- [9] L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. Xing, et al., "Judging Ilm-as-a-judge with mt-bench and chatbot arena," Advances in Neural Information Processing Systems, vol. 36, pp. 46595–46623, 2023.
- [10] H. Touvron, T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro, F. Azhar, et al., "Llama: Open and efficient foundation language models," arXiv preprint arXiv:2302.13971, 2023.
- [11] A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Yang, A. Fan, et al., "The llama 3 herd of models," arXiv preprint arXiv:2407.21783, 2024.
- [12] A. Yang, B. Yang, B. Hui, B. Zheng, B. Yu, C. Zhou, C. Li, C. Li, D. Liu, F. Huang, G. Dong, H. Wei, H. Lin, J. Tang, J. Wang, J. Yang, J. Tu, J. Zhang, J. Ma, J. Yang, J. Xu, J. Zhou, J. Bai, J. He, J. Lin, K. Dang, K. Lu, K. Chen, K. Yang, M. Li, M. Xue, N. Ni, P. Zhang, P. Wang, R. Peng, R. Men, R. Gao, R. Lin, S. Wang, S. Bai, S. Tan, T. Zhu, T. Li, T. Liu, W. Ge, X. Deng, X. Zhou, X. Ren, X. Zhang, X. Wei, X. Ren, X. Liu, Y. Fan, Y. Yao, Y. Zhang, Y. Wan, Y. Chu, Y. Liu, Z. Cui, Z. Zhang, Z. Guo, and Z. Fan, "Qwen2 technical report," 2024.
- [13] A. Yang, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Li, D. Liu, F. Huang, H. Wei, et al., "Qwen2. 5 technical report," arXiv preprint arXiv:2412.15115, 2024.
- [14] J. Wei, M. Bosma, V. Zhao, K. Guu, A. W. Yu, B. Lester, N. Du, A. M. Dai, and Q. V. Le, "Finetuned language models are zero-shot learners," in *International Conference on Learning Representations*, 2022.
- [15] H. W. Chung, L. Hou, S. Longpre, B. Zoph, Y. Tay, W. Fedus, Y. Li, X. Wang, M. Dehghani, S. Brahma, *et al.*, "Scaling instruction-finetuned language models," *arXiv preprint arXiv:2210.11416*, 2022.
- [16] P. Wang, S. Bai, S. Tan, S. Wang, Z. Fan, J. Bai, K. Chen, X. Liu, J. Wang, W. Ge, Y. Fan, K. Dang, M. Du, X. Ren, R. Men, D. Liu, C. Zhou, J. Zhou, and J. Lin, "Qwen2-vl: Enhancing vision-language model's perception of the world at any resolution," 2024.

- [17] Z. Chen, W. Wang, Y. Cao, Y. Liu, Z. Gao, E. Cui, J. Zhu, S. Ye, H. Tian, Z. Liu, *et al.*, "Expanding performance boundaries of open-source multimodal models with model, data, and test-time scaling," *arXiv preprint arXiv:2412.05271*, 2024.
- [18] J. Zhu, W. Wang, Z. Chen, Z. Liu, S. Ye, L. Gu, Y. Duan, H. Tian, W. Su, J. Shao, et al., "Internvl3: Exploring advanced training and test-time recipes for open-source multimodal models," arXiv preprint arXiv:2504.10479, 2025.
- [19] A. Masry, D. X. Long, J. Q. Tan, S. Joty, and E. Hoque, "Chartqa: A benchmark for question answering about charts with visual and logical reasoning," arXiv preprint arXiv:2203.10244, 2022.
- [20] Y. Liu, H. Duan, Y. Zhang, B. Li, S. Zhang, W. Zhao, Y. Yuan, J. Wang, C. He, Z. Liu, et al., "Mmbench: Is your multi-modal model an all-around player?," arXiv preprint arXiv:2307.06281, 2023.
- [21] X. Yue, T. Zheng, Y. Ni, Y. Wang, K. Zhang, S. Tong, Y. Sun, B. Yu, G. Zhang, H. Sun, et al., "Mmmu-pro: A more robust multi-discipline multimodal understanding benchmark," arXiv preprint arXiv:2409.02813, 2024.
- [22] L. Chen, J. Li, X. Dong, P. Zhang, Y. Zang, Z. Chen, H. Duan, J. Wang, Y. Qiao, D. Lin, *et al.*, "Are we on the right way for evaluating large vision-language models?," *arXiv preprint arXiv:2403.20330*, 2024.
- [23] B. Li, R. Wang, G. Wang, Y. Ge, Y. Ge, and Y. Shan, "Seed-bench: Benchmarking multimodal llms with generative comprehension," *arXiv preprint arXiv:2307.16125*, 2023.
- [24] B. Li, Y. Ge, Y. Chen, Y. Ge, R. Zhang, and Y. Shan, "Seed-bench-2-plus: Benchmarking multimodal large language models with text-rich visual comprehension," *arXiv preprint arXiv:2404.16790*, 2024.
- [25] J. Kaplan, S. McCandlish, T. Henighan, T. B. Brown, B. Chess, R. Child, S. Gray, A. Radford, J. Wu, and D. Amodei, "Scaling laws for neural language models," *arXiv preprint arXiv:2001.08361*, 2020.
- [26] A. Chowdhery, S. Narang, J. Devlin, M. Bosma, G. Mishra, A. Roberts, P. Barham, H. W. Chung, C. Sutton, S. Gehrmann, *et al.*, "Palm: Scaling language modeling with pathways," *Journal of Machine Learning Research*, vol. 24, no. 240, pp. 1–113, 2023.
- [27] Y. Li, Y. Zhang, C. Wang, Z. Zhong, Y. Chen, R. Chu, S. Liu, and J. Jia, "Mini-gemini: Mining the potential of multi-modality vision language models," *arXiv preprint arXiv:2403.18814*, 2024.
- [28] B. Li, Y. Zhang, D. Guo, R. Zhang, F. Li, H. Zhang, K. Zhang, Y. Li, Z. Liu, and C. Li, "Llava-onevision: Easy visual task transfer," *arXiv preprint arXiv:2408.03326*, 2024.
- [29] X. Yue, T. Zheng, G. Zhang, and W. Chen, "Mammoth2: Scaling instructions from the web," 2024.
- [30] S. Gu, J. Zhang, S. Zhou, K. Yu, Z. Xing, L. Wang, Z. Cao, J. Jia, Z. Zhang, Y. Wang, *et al.*, "Infinity-mm: Scaling multimodal performance with large-scale and high-quality instruction data," *arXiv preprint arXiv:2410.18558*, 2024.
- [31] B.-K. Lee, S. Chung, C. W. Kim, B. Park, and Y. M. Ro, "Trol: Traversal of layers for large language and vision models," *arXiv preprint arXiv:2406.12246*, 2024.
- [32] B.-K. Lee, S. Chung, C. W. Kim, B. Park, and Y. M. Ro, "Phantom of latent for large language and vision models," *arXiv preprint arXiv:2409.14713*, 2024.
- [33] B.-K. Lee, B. Park, C. W. Kim, and Y. M. Ro, "Collavo: Crayon large language and vision model," *arXiv preprint arXiv:2402.11248*, 2024.
- [34] B.-K. Lee, B. Park, C. W. Kim, and Y. M. Ro, "Moai: Mixture of all intelligence for large language and vision models," *arXiv preprint arXiv:2403.07508*, 2024.

- [35] B.-K. Lee, C. W. Kim, B. Park, and Y. M. Ro, "Meteor: Mamba-based traversal of rationale for large language and vision models," *arXiv preprint arXiv:2405.15574*, 2024.
- [36] Q. Yu, Z. Zhang, R. Zhu, Y. Yuan, X. Zuo, Y. Yue, T. Fan, G. Liu, L. Liu, X. Liu, et al., "Dapo: An open-source llm reinforcement learning system at scale," arXiv preprint arXiv:2503.14476, 2025.
- [37] W. Huang, B. Jia, Z. Zhai, S. Cao, Z. Ye, F. Zhao, Z. Xu, Y. Hu, and S. Lin, "Vision-r1: Incentivizing reasoning capability in multimodal large language models," *arXiv preprint arXiv:2503.06749*, 2025.
- [38] Z. Lu, Y. Chai, Y. Guo, X. Yin, L. Liu, H. Wang, G. Xiong, and H. Li, "Ui-r1: Enhancing action prediction of gui agents by reinforcement learning," *arXiv preprint arXiv:2503.21620*, 2025.
- [39] E. Yu, K. Lin, L. Zhao, J. Yin, Y. Wei, Y. Peng, H. Wei, J. Sun, C. Han, Z. Ge, *et al.*, "Perception-r1: Pioneering perception policy with reinforcement learning," *arXiv preprint arXiv:2504.07954*, 2025.
- [40] X. Liu, J. Ni, Z. Wu, C. Du, L. Dou, H. Wang, T. Pang, and M. Q. Shieh, "Noisyrollout: Reinforcing visual reasoning with data augmentation," *arXiv preprint arXiv:2504.13055*, 2025.
- [41] H. Shen, P. Liu, J. Li, C. Fang, Y. Ma, J. Liao, Q. Shen, Z. Zhang, K. Zhao, Q. Zhang, et al., "Vlm-r1: A stable and generalizable r1-style large vision-language model," arXiv preprint arXiv:2504.07615, 2025.
- [42] L. Chen, L. Li, H. Zhao, Y. Song, and Vinci, "R1-v: Reinforcing super generalization ability in vision-language models with less than \$3." https://github.com/Deep-Agent/R1-V, 2025. Accessed: 2025-02-02.
- [43] Z. Shao, P. Wang, Q. Zhu, R. Xu, J. Song, X. Bi, H. Zhang, M. Zhang, Y. Li, Y. Wu, *et al.*, "Deepseekmath: Pushing the limits of mathematical reasoning in open language models," *arXiv* preprint arXiv:2402.03300, 2024.
- [44] J. Ho and S. Ermon, "Generative adversarial imitation learning," *Advances in neural information processing systems*, vol. 29, 2016.
- [45] B.-K. Lee, R. Hachiuma, Y.-C. F. Wang, Y. M. Ro, and Y.-H. Wu, "Vlsi: Verbalized layers-to-interactions from large to small vision language models," *arXiv preprint arXiv:2412.01822*, 2024.
- [46] Y. Li, F. Wei, C. Zhang, and H. Zhang, "Eagle: Speculative sampling requires rethinking feature uncertainty," *arXiv preprint arXiv:2401.15077*, 2024.
- [47] T. Salimans, I. Goodfellow, W. Zaremba, V. Cheung, A. Radford, and X. Chen, "Improved techniques for training gans," *Advances in neural information processing systems*, vol. 29, 2016.
- [48] M. Heusel, H. Ramsauer, T. Unterthiner, B. Nessler, and S. Hochreiter, "Gans trained by a two time-scale update rule converge to a local nash equilibrium," *Advances in neural information processing systems*, vol. 30, 2017.
- [49] Z. Liu, Z. Sun, Y. Zang, X. Dong, Y. Cao, H. Duan, D. Lin, and J. Wang, "Visual-rft: Visual reinforcement fine-tuning," *arXiv preprint arXiv:2503.01785*, 2025.
- [50] A. Hurst, A. Lerer, A. P. Goucher, A. Perelman, A. Ramesh, A. Clark, A. Ostrow, A. Welihinda, A. Hayes, A. Radford, et al., "Gpt-4o system card," arXiv preprint arXiv:2410.21276, 2024.
- [51] G. Team, R. Anil, S. Borgeaud, Y. Wu, J.-B. Alayrac, J. Yu, R. Soricut, J. Schalkwyk, A. M. Dai, A. Hauth, *et al.*, "Gemini: a family of highly capable multimodal models," *arXiv preprint arXiv:2312.11805*, 2023.
- [52] Anthropic, "The claude 3 model family: Opus, sonnet, haiku." https://www.anthropic.com, 2024.

- [53] M. Deitke, C. Clark, S. Lee, R. Tripathi, Y. Yang, J. S. Park, M. Salehi, N. Muennighoff, K. Lo, L. Soldaini, *et al.*, "Molmo and pixmo: Open weights and open data for state-of-the-art multimodal models," *arXiv preprint arXiv:2409.17146*, 2024.
- [54] W. Dai, N. Lee, B. Wang, Z. Yang, Z. Liu, J. Barker, T. Rintamaki, M. Shoeybi, B. Catanzaro, and W. Ping, "Nvlm: Open frontier-class multimodal llms," *arXiv preprint*, 2024.
- [55] Z. Zong, B. Ma, D. Shen, G. Song, H. Shao, D. Jiang, H. Li, and Y. Liu, "Mova: Adapting mixture of vision experts to multimodal context," *arXiv preprint arXiv:2404.13046*, 2024.
- [56] M. Shi, F. Liu, S. Wang, S. Liao, S. Radhakrishnan, D.-A. Huang, H. Yin, K. Sapra, Y. Yacoob, H. Shi, et al., "Eagle: Exploring the design space for multimodal llms with mixture of encoders," arXiv preprint arXiv:2408.15998, 2024.
- [57] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, "Learning transferable visual models from natural language supervision," in *Proceedings of the 38th International Conference on Machine Learning* (M. Meila and T. Zhang, eds.), vol. 139 of *Proceedings of Machine Learning Research*, pp. 8748–8763, PMLR, 18–24 Jul 2021.
- [58] S. Woo, S. Debnath, R. Hu, X. Chen, Z. Liu, I. S. Kweon, and S. Xie, "Convnext v2: Codesigning and scaling convnets with masked autoencoders," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 16133–16142, 2023.
- [59] M. Oquab, T. Darcet, T. Moutakanni, H. Vo, M. Szafraniec, V. Khalidov, P. Fernandez, D. Haziza, F. Massa, A. El-Nouby, *et al.*, "Dinov2: Learning robust visual features without supervision," *arXiv preprint arXiv:2304.07193*, 2023.
- [60] A. Kirillov, E. Mintun, N. Ravi, H. Mao, C. Rolland, L. Gustafson, T. Xiao, S. Whitehead, A. C. Berg, W.-Y. Lo, et al., "Segment anything," in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 4015–4026, 2023.
- [61] Y. Peng, G. Zhang, M. Zhang, Z. You, J. Liu, Q. Zhu, K. Yang, X. Xu, X. Geng, and X. Yang, "Lmm-r1: Empowering 3b lmms with strong reasoning abilities through two-stage rule-based rl," arXiv preprint arXiv:2503.07536, 2025.
- [62] Y. Deng, H. Bansal, F. Yin, N. Peng, W. Wang, and K.-W. Chang, "Openvlthinker: An early exploration to complex vision-language reasoning via iterative self-improvement," *arXiv* preprint arXiv:2503.17352, 2025.
- [63] Y. Yang, X. He, H. Pan, X. Jiang, Y. Deng, X. Yang, H. Lu, D. Yin, F. Rao, M. Zhu, et al., "R1-onevision: Advancing generalized multimodal reasoning through cross-modal formalization," arXiv preprint arXiv:2503.10615, 2025.
- [64] H. Zhou, X. Li, R. Wang, M. Cheng, T. Zhou, and C.-J. Hsieh, "R1-zero's" aha moment" in visual reasoning on a 2b non-sft model," *arXiv preprint arXiv:2503.05132*, 2025.
- [65] F. Meng, L. Du, Z. Liu, Z. Zhou, Q. Lu, D. Fu, T. Han, B. Shi, W. Wang, J. He, *et al.*, "Mm-eureka: Exploring the frontiers of multimodal reasoning with rule-based reinforcement learning," *arXiv preprint arXiv:2503.07365*, 2025.
- [66] A. Hussein, M. M. Gaber, E. Elyan, and C. Jayne, "Imitation learning: A survey of learning methods," *ACM Computing Surveys (CSUR)*, vol. 50, no. 2, pp. 1–35, 2017.
- [67] W. Yu, Z. Yang, L. Ren, L. Li, J. Wang, K. Lin, C.-C. Lin, Z. Liu, L. Wang, and X. Wang, "Mm-vet v2: A challenging benchmark to evaluate large multimodal models for integrated capabilities," *arXiv preprint arXiv:2408.00765*, 2024.
- [68] O. Contributors, "Opencompass: A universal evaluation platform for foundation models." https://github.com/open-compass/opencompass, 2023.
- [69] Z. Chen, J. Wu, W. Wang, W. Su, G. Chen, S. Xing, Z. Muyan, Q. Zhang, X. Zhu, L. Lu, *et al.*, "Internvl: Scaling up vision foundation models and aligning for generic visual-linguistic tasks," *arXiv* preprint arXiv:2312.14238, 2023.

- [70] Y. Yao, T. Yu, A. Zhang, C. Wang, J. Cui, H. Zhu, T. Cai, H. Li, W. Zhao, Z. He, et al., "Minicpm-v: A gpt-4v level mllm on your phone," arXiv preprint arXiv:2408.01800, 2024.
- [71] S. Lu, Y. Li, Q.-G. Chen, Z. Xu, W. Luo, K. Zhang, and H.-J. Ye, "Ovis: Structural embedding alignment for multimodal large language model," *arXiv:2405.20797*, 2024.
- [72] M. Abdin, S. A. Jacobs, A. A. Awan, J. Aneja, A. Awadallah, H. Awadalla, N. Bach, A. Bahree, A. Bakhtiari, H. Behl, *et al.*, "Phi-3 technical report: A highly capable language model locally on your phone," *arXiv preprint arXiv:2404.14219*, 2024.
- [73] A. Abouelenin, A. Ashfaq, A. Atkinson, H. Awadalla, N. Bach, J. Bao, A. Benhaim, M. Cai, V. Chaudhary, C. Chen, *et al.*, "Phi-4-mini technical report: Compact yet powerful multimodal language models via mixture-of-loras," *arXiv preprint arXiv:2503.01743*, 2025.
- [74] OpenAI, "Gpt-4v(ision) system card," 2023. https://openai.com/research/gpt-4v-system-card, Last accessed on 2024-02-13.
- [75] Y. Gu, L. Dong, F. Wei, and M. Huang, "Minillm: Knowledge distillation of large language models," in *The Twelfth International Conference on Learning Representations*, 2024.
- [76] J. Ko, S. Kim, T. Chen, and S.-Y. Yun, "Distillm: Towards streamlined distillation for large language models," *arXiv preprint arXiv:2402.03898*, 2024.
- [77] Y. Cai, J. Zhang, H. He, X. He, A. Tong, Z. Gan, C. Wang, and X. Bai, "Llava-kd: A framework of distilling multimodal large language models," *arXiv preprint arXiv:2410.16236*, 2024.
- [78] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," *arXiv* preprint arXiv:1503.02531, 2015.
- [79] W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. Gonzalez, H. Zhang, and I. Stoica, "Efficient memory management for large language model serving with pagedattention," in *Proceedings of the 29th Symposium on Operating Systems Principles*, pp. 611–626, 2023.
- [80] S. Rajbhandari, J. Rasley, O. Ruwase, and Y. He, "Zero: Memory optimizations toward training trillion parameter models," in *SC20: International Conference for High Performance Computing, Networking, Storage and Analysis*, pp. 1–16, IEEE, 2020.
- [81] I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *International Conference on Learning Representations*, 2019.
- [82] Z. Liu, L. Zhu, B. Shi, Z. Zhang, Y. Lou, S. Yang, H. Xi, S. Cao, Y. Gu, D. Li, et al., "Nvila: Efficient frontier visual language models," arXiv preprint arXiv:2412.04468, 2024.
- [83] S. Singh, A. Yadav, J. Jain, H. Shi, J. Johnson, and K. Desai, "Benchmarking object detectors with coco: A new path forward," in *European Conference on Computer Vision*, pp. 279–295, Springer, 2024.
- [84] G. Van Horn, O. Mac Aodha, Y. Song, Y. Cui, C. Sun, A. Shepard, H. Adam, P. Perona, and S. Belongie, "The inaturalist species classification and detection dataset," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 8769–8778, 2018.
- [85] Y. Goyal, T. Khot, D. Summers-Stay, D. Batra, and D. Parikh, "Making the V in VQA matter: Elevating the role of image understanding in Visual Question Answering," in *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.
- [86] Z. Li, X. Wang, E. Stengel-Eskin, A. Kortylewski, W. Ma, B. Van Durme, and A. L. Yuille, "Super-clevr: A virtual benchmark to diagnose domain robustness in visual reasoning," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 14963–14973, 2023.
- [87] R. Zhang, X. Wei, D. Jiang, Y. Zhang, Z. Guo, C. Tong, J. Liu, A. Zhou, B. Wei, S. Zhang, et al., "Mavis: Mathematical visual instruction tuning," arXiv e-prints, pp. arXiv–2407, 2024.

- [88] P. Lu, R. Gong, S. Jiang, L. Qiu, S. Huang, X. Liang, and S.-C. Zhu, "Inter-gps: Interpretable geometry problem solving with formal language and symbolic reasoning," in *The 59th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2021.
- [89] P. Lu, S. Mishra, T. Xia, L. Qiu, K.-W. Chang, S.-C. Zhu, O. Tafjord, P. Clark, and A. Kalyan, "Learn to explain: Multimodal reasoning via thought chains for science question answering," *Advances in Neural Information Processing Systems*, vol. 35, pp. 2507–2521, 2022.
- [90] Y. Zhang, R. Zhang, J. Gu, Y. Zhou, N. Lipka, D. Yang, and T. Sun, "Llavar: Enhanced visual instruction tuning for text-rich image understanding," arXiv preprint arXiv:2306.17107, 2023.
- [91] F. Liu, G. Emerson, and N. Collier, "Visual spatial reasoning," *Transactions of the Association for Computational Linguistics*, vol. 11, pp. 635–651, 2023.
- [92] M. Acharya, K. Kafle, and C. Kanan, "Tallyqa: Answering complex counting questions," in *Proceedings of the AAAI conference on artificial intelligence*, vol. 33, pp. 8076–8084, 2019.
- [93] P. Lu, L. Qiu, K.-W. Chang, Y. N. Wu, S.-C. Zhu, T. Rajpurohit, P. Clark, and A. Kalyan, "Dynamic prompt learning via policy gradient for semi-structured mathematical reasoning," arXiv preprint arXiv:2209.14610, 2022.
- [94] V. Hosu, H. Lin, T. Sziranyi, and D. Saupe, "Koniq-10k: An ecologically valid database for deep learning of blind image quality assessment," *IEEE Transactions on Image Processing*, vol. 29, pp. 4041–4056, 2020.
- [95] W. Wang, Z. Chen, W. Wang, Y. Cao, Y. Liu, Z. Gao, J. Zhu, X. Zhu, L. Lu, Y. Qiao, *et al.*, "Enhancing the reasoning ability of multimodal large language models via mixed preference optimization," *arXiv preprint arXiv:2411.10442*, 2024.
- [96] T. Yu, H. Zhang, Y. Yao, Y. Dang, D. Chen, X. Lu, G. Cui, T. He, Z. Liu, T.-S. Chua, *et al.*, "Rlaif-v: Aligning mllms through open-source ai feedback for super gpt-4v trustworthiness," *arXiv* preprint arXiv:2405.17220, 2024.
- [97] A. D. Lindström and S. S. Abraham, "Clevr-math: A dataset for compositional language, visual and mathematical reasoning," *arXiv preprint arXiv:2208.05358*, 2022.
- [98] Z. Huang, K. Chen, J. He, X. Bai, D. Karatzas, S. Lu, and C. Jawahar, "Icdar2019 competition on scanned receipt ocr and information extraction," in *2019 International Conference on Document Analysis and Recognition (ICDAR)*, pp. 1516–1520, IEEE, 2019.
- [99] M. Mathew, D. Karatzas, and C. Jawahar, "Docvqa: A dataset for vqa on document images," in *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, pp. 2200–2209, 2021.
- [100] S. E. Kahou, V. Michalski, A. Atkinson, Á. Kádár, A. Trischler, and Y. Bengio, "Figureqa: An annotated figure dataset for visual reasoning," *arXiv preprint arXiv:1710.07300*, 2017.
- [101] D. A. Hudson and C. D. Manning, "Gqa: A new dataset for real-world visual reasoning and compositional question answering," in *Proceedings of the IEEE/CVF conference on computer* vision and pattern recognition, pp. 6700–6709, 2019.
- [102] M. Mathew, V. Bagal, R. Tito, D. Karatzas, E. Valveny, and C. Jawahar, "Infographicvqa," in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pp. 1697–1706, 2022.
- [103] Q. Chen, L. Qin, J. Zhang, Z. Chen, X. Xu, and W. Che, "M³cot: A novel benchmark for multi-domain multi-step multi-modal chain-of-thought," in *Proc. of ACL*, 2024.
- [104] S. Chang, D. Palzer, J. Li, E. Fosler-Lussier, and N. Xiao, "Mapqa: A dataset for question answering on choropleth maps," *arXiv preprint arXiv:2211.08545*, 2022.
- [105] K. Marino, M. Rastegari, A. Farhadi, and R. Mottaghi, "Ok-vqa: A visual question answering benchmark requiring external knowledge," in *Proceedings of the IEEE/cvf conference on computer vision and pattern recognition*, pp. 3195–3204, 2019.

- [106] A. Singh, V. Natarajan, M. Shah, Y. Jiang, X. Chen, D. Batra, D. Parikh, and M. Rohrbach, "Towards vqa models that can read," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 8317–8326, 2019.
- [107] Y. Lu, D. Jiang, W. Chen, W. Y. Wang, Y. Choi, and B. Y. Lin, "Wildvision: Evaluating vision-language models in the wild with human preferences," arXiv preprint arXiv:2406.11069, 2024.
- [108] K. Kafle, B. Price, S. Cohen, and C. Kanan, "Dvqa: Understanding data visualizations via question answering," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 5648–5656, 2018.
- [109] J. Cao and J. Xiao, "An augmented benchmark dataset for geometric question answering through dual parallel text encoding," in *Proceedings of the 29th international conference on computational linguistics*, pp. 1511–1520, 2022.
- [110] M. Seo, H. Hajishirzi, A. Farhadi, O. Etzioni, and C. Malcolm, "Solving geometry problems: Combining text and diagram interpretation," in *Proceedings of the 2015 conference on empirical methods in natural language processing*, pp. 1466–1476, 2015.
- [111] P. Lu, L. Qiu, J. Chen, T. Xia, Y. Zhao, W. Zhang, Z. Yu, X. Liang, and S.-C. Zhu, "Iconqa: A new benchmark for abstract diagram understanding and visual language reasoning," *arXiv* preprint arXiv:2110.13214, 2021.
- [112] J. Chen, T. Li, J. Qin, P. Lu, L. Lin, C. Chen, and X. Liang, "Unigeo: Unifying geometry logical reasoning via reformulating mathematical expression," *arXiv preprint arXiv:2212.02746*, 2022.
- [113] M. Kazemi, H. Alvari, A. Anand, J. Wu, X. Chen, and R. Soricut, "Geomverse: A systematic evaluation of large models for geometric reasoning," *arXiv preprint arXiv:2312.12241*, 2023.
- [114] J. Gao, R. Pi, J. Zhang, J. Ye, W. Zhong, Y. Wang, L. Hong, J. Han, H. Xu, Z. Li, et al., "G-llava: Solving geometric problem with multi-modal large language model," arXiv preprint arXiv:2312.11370, 2023.
- [115] W. Shi, Z. Hu, Y. Bin, J. Liu, Y. Yang, S.-K. Ng, L. Bing, and R. K.-W. Lee, "Math-llava: Bootstrapping mathematical reasoning for multimodal large language models," *arXiv preprint arXiv:2406.17294*, 2024.
- [116] e. a. Yunxin Li, "Cognitive visual-language mapper: Advancing multimodal comprehension with enhanced visual knowledge alignment," *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.
- [117] Y. Chen, H. Hu, Y. Luan, H. Sun, S. Changpinyo, A. Ritter, and M.-W. Chang, "Can pre-trained vision and language models answer visual information-seeking questions?," *arXiv preprint arXiv:2302.11713*, 2023.
- [118] X. Huang, Y.-J. Huang, Y. Zhang, W. Tian, R. Feng, Y. Zhang, Y. Xie, Y. Li, and L. Zhang, "Open-set image tagging with multi-grained text supervision," arXiv preprint arXiv:2310.15200, 2023.
- [119] B.-K. Lee, J. Kim, and Y. M. Ro, "Masking adversarial damage: Finding adversarial saliency for robust and sparse network," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 15126–15136, 2022.
- [120] J. Kim, B.-K. Lee, and Y. M. Ro, "Distilling robust and non-robust features in adversarial examples by information bottleneck," *Advances in Neural Information Processing Systems*, vol. 34, pp. 17148–17159, 2021.
- [121] J. Kim, B.-K. Lee, and Y. M. Ro, "Demystifying causal features on adversarial examples and causal inoculation for robust network by adversarial instrumental variable regression," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 12302–12312, 2023.

- [122] B.-K. Lee, Y. Yu, and Y. M. Ro, "Towards adversarial robustness of bayesian neural network through hierarchical variational inference," 2021.
- [123] B.-K. Lee, R. Hachiuma, Y. M. Ro, Y.-C. F. Wang, and Y.-H. Wu, "Genrecal: Generation after recalibration from large to small vision-language models," *arXiv preprint arXiv:2506.15681*, 2025.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: Abstract and Introduction section, and Figure 1-2

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Discussion and Limitation section

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [NA]

Justification: There are no theoretical formulations that need to be proved.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and crossreferenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Experiment section and Appendix C

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
- (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
- (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [NA]

Justification: Dataset is all open-source, but the code and the model checkpoints will be open once it is accepted.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (https://nips.cc/public/guides/CodeSubmissionPolicy) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so "No" is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (https://nips.cc/public/guides/CodeSubmissionPolicy) for more details.
- The authors should provide instructions on data access and preparation, including how
 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: Experiment section and Appendix C

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [NA]
Justification: N/A

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).

- It should be clear whether the error bar is the standard deviation or the standard error
 of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Experiment section and Appendix C

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics https://neurips.cc/public/EthicsGuidelines?

Answer: [Yes]

Justification: We have read all materials regarding the NeurIPS code of ethics.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Appendix D

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]
Justification: N/A

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Method, Experiment section, Appendix C

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.

• If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: Experiment Section and Appendix C

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]
Justification: N/A
Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]
Justification: N/A
Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [NA]
Justification: N/A

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (https://neurips.cc/Conferences/2025/LLM) for what should or should not be described.

A Algorithm of RL only with Answer Rewards from LLM-as-a-Judge

Algorithm 1 RL purely with GRPO or Dr.GRPO based on accuracy rewards from LLM-as-a-Judge

```
Require: Pre-trained VLMs \pi_{\theta_{\text{init}}}

1: Set reference model \pi_{\text{ref}} \leftarrow \pi_{\theta_{\text{init}}}

2: Set the training model \pi_{\theta} \leftarrow \pi_{\theta_{\text{init}}}

3: for sample a batch \mathcal{B} in Dataset do

4: Copy and freeze model \pi_{\theta_{\text{old}}} \leftarrow \pi_{\theta}

5: Sample G outputs \{o_i\}_{i=1}^G \sim \pi_{\theta_{\text{old}}}(\cdot|q) for each question q \in \mathcal{B}

6: Compute Rewards and Advantages for all G outputs by LLM-as-a-Judge

7: for VLMs Updating iteration = 1, 2, \dots, \mu do

8: Update \pi_{\theta} by using the objective of GRPO [43] or Dr.GRPO [8]

9: end for
```

B Algorithm of RIL for VLMs

```
Algorithm 2 RIL for VLMs
Require: Pre-trained discriminator D_{\phi} and Pre-trained student VLMs \pi_{\theta_{\text{min}}}
Require: Saved collection for the generated text responses \{o^{(1)}\} from teacher VLMs
 1: Set reference model \pi_{\text{ref}} \leftarrow \pi_{\theta_{\text{init}}}
 2: Set the training model \pi_{\theta} \leftarrow \pi_{\theta_{\text{init}}}
 3: for sample a batch \mathcal{B} in Dataset do
        Copy and freeze model \pi_{\theta_{\text{old}}} \leftarrow \pi_{\theta}
        Sample G outputs \{o_i^{(\mathrm{s})}\}_{i=1}^G \sim \pi_{\theta_{\mathrm{old}}}(\cdot|q) for each question q \in \mathcal{B}
 5:
        Extract G outputs \{o_i^{(t)}\}_{i=1}^G in the saved collection for each question q \in \mathcal{B}
 6:
 7:
        for Discriminator Updating iteration = 1, 2, \dots, \mu do
 8:
            Update D_{\phi} by using Equation 1
 9:
        end for
        Compute Rewards and Advantages for all 2G outputs by Discriminator and LLM-as-a-Judge
10:
11:
        for Student VLMs Updating iteration = 1, 2, \dots, \mu do
            Update \pi_{\theta} by using Equation 2
12:
13:
        end for
14: end for
```

C Implementation Detail

Details of Training and Evaluation. We train and evaluate RIL, mainly on NVIDIA A100 80GB GPUs. Because it is practically critical point to get fast text generation during training, we utilize vLLM [79] built on PagedAttention. For the pre-training step of discriminator, we assign vLLM [79] to 8 GPUs for fast text generation, and we generate N=16 text responses for each question at student and teacher VLMs, respectively. Once it is finished, we use DeepSpeed engine with ZeRO-3 [80] for 8 GPUs, and we use AdamW optimizer [81] and apply a linearly decayed learning rate from 1e-5 to 1e-6 to pre-training discriminator and SFT of student VLMs. In subsequent step, when training both discriminator and student VLMs, one of 8 GPUs is assigned to conduct online generation of student VLMs by vLLM [79]. In addition, another GPU is assigned to conduct LLM-as-a-Judge [9] by vLLM [79]. The other 6 GPUs are assigned to use DeepSpeed engine with ZeRO-3 [80]. Mimicking step requires static learning rate 1e-6 and μ =1 iteration to train both student VLMs and the discriminator. Note that, when we generate text responses, we generate G=4 responses for each question, by setting temperature to 1.0, top-p to 0.95, top-k to 50, and repetition penalty to 1.05, in order to get diverse text responses. For stable training, we handle large batch sizes by using gradient accumulation with 6 steps. At every step, we use 4 batches per one GPU, leading to total 144 batches. For evaluation, we use distributed data parallel to load student VLMs for all 8 GPUs and

score the correctness or generation quality of their text responses by using their default generation hyperparameter.

Computational Cost. Compared to GRPO [43], RIL requires increased computational costs during training due to teacher, student, discriminator, and LLM-as-a-Judge [9]. However, we would like to emphasize that discriminator and student costs can be at least mitigated. While pretraining the discriminator is a necessary step and cannot be avoided from a computational standpoint, we can reduce training costs of discriminator and student during the RIL loop (Algorithm 2) by managing model weights efficiently between CPU and GPU. Specifically, we offload the weights of the discriminator and student to the CPU when they are not in use. When needed, we load the appropriate weights from the CPU to the GPU using DeepSpeed API [80]. This allows us to load or unload model weights without uploading both architecture to GPU, thereby minimizing computational overhead. This approach is feasible because we use the same model architecture for both the discriminator and student, enabling seamless weight replacement and resource-efficient training.

Visual Instruction Tuning Dataset. We assemble a 4M-sample visual instruction tuning dataset that encompasses a diverse array of vision-language tasks, including general visual question answering, dense image captioning, chart, diagram, and document understanding, common-sense knowledge, scientific and mathematical problem-solving, and multidimensional reasoning. For SFT, we utilize the entire 4M-sample dataset, and then we curate a 40K-sample dataset for RIL of VLMs based on log-probability sampling [82] and overlong filtering [36]. Our dataset integrates both real-world and synthetic sources: COCO-ReM [83], iNaturalist2018 [84], VQA-v2 [85], Super-CLEVR [86], MAVIS [87], Geometry3K [88], SQA [89], AI2D [2], SA-1B [60], LLaVAR [90], VSR [91], TallyQA [92], TabMWP [93], KonIQ [94], InternVL [95]-filtered synthetic knowledge dataset covering politics, math, physics, chemistry, RLAI-F [96], CLEVR-Math [97], SROIE [98], ChartQA [19], DocVQA [99], FigureQA [100], GQA [101], InfoVQA [102], M3CoT [103], MapQA [104], OK-VQA [105], TextVQA [106], WildVision [107], DVQA [108], GeoQA+ [109], GeOS [110], IconQA [111], UniGEO [112], GeomVerse [113], Geo170K [114], MathV360K [115], multimodal wikipedia knowledge [116], InfoSeek [117], and RAM++ [118]-filtered synthetic data of Infinity-MM [30] covering coarse and fine-grained perception, relation, attribute, and logic reasoning.

D Broader Impacts

Our reinforcement and imitation learning (RIL) approach democratises access to advanced vision—language capabilities by enabling lightweight, efficient models to run on edge devices and in low-resource settings—thereby lowering both deployment costs and environmental impact. RIL fosters seamless interoperability across diverse model architectures and tokenization schemes. The use of multiple large teacher VLMs to generate varied training samples further accelerates performance gains, underscoring the value of collaborative model development. We believe RIL will inspire future research into efficient multimodal learning techniques and support the widespread, sustainable adoption of lightweight [119] and robust [120–122] AI solutions across a broad range of applications even with different tokenizers [123]. We hope that by bridging efficiency and performance, RIL become a landscape of building sophisticated multimodal AI technologies and accelerating innovation across industries and academia.